

Model Latency

* Includes network latency as well.

- model: qwen2:72b-instruct-q6_K

```
Using model: qwen2:72b-instruct-q6_K
Sample count: 300
Applying Classification: 100% ██████████ | 300/300 [28:57<00:00, 5.79s/it]
Evaluating... ⌚
Evaluation completed. Check the output files (qwen2_72b-instruct-q6_K).
```

Average : 5.8 Sec

- llama3.3:70b-instruct-q6_K

```
Using model: llama3.3:70b-instruct-q6_K
Sample count: 300
Applying Classification: 100% | 300/300 [18:05<00:00, 3.62s/it]
```

Average : 3.60 Sec

- gemma2:27b-instruct-q6_K

```
Using model: gemma2:27b-instruct-q6_K
Sample count: 300
Applying Classification: 100% | 300/300 [02:34<00:00, 1.94it/s]
```

Average : < 1 Sec

- gemma:7b-instruct-q6_K

```
Using model: gemma:7b-instruct-q6_K
Sample count: 300
Applying Classification: 100% | 300/300 [07:18<00:00, 1.46s/it]
```

Average : 1.46 Sec

- phi3:14b-medium-128k-instruct-q6_K

```
Using model: phi3:14b-medium-128k-instruct-q6_K
Sample count: 300
Applying Classification: 100% | 300/300 [04:46<00:00, 1.05it/s]
```

Average : < 1sec

- llama3.1:70b-instruct-q6_K

```
Using model: llama3.1:70b-instruct-q6_K
Sample count: 300
Applying Classification: 100% | 300/300 [17:29<00:00, 3.50s/it]
```

Average : 3.50 Sec